



database/sql

There and back and again


Intro

- Huu Khiem (Mark), currently at Viki (viki.com)
- Motivation for talk: need to understand how Go + SQL works better, to:
 - Optimize code (medium/large services PostgreSQL, 8M+ users, lots of touch point with DB)
 - Capacity planning (understand behavior at average & peak load)
 - Failure handling (know when and how database code fails)
 - Write less (lots of boilerplate code to delete)
- Talk based on Go 1.7. Lots of change in Go 1.8!


It's harder than it looks

```
db, err := sql.Open("postgres", "...")
if err != nil {
    // handle error
}


// ...
rows, err := db.Query(query, args...)
if err != nil {
    // handle error
}
defer row.Close()
for rows.Next() {
    // scan and parse
}
```




Talking to PostgreSQL, high level



Establish connection



Establish connection (cont.)



Tracing the calls: serialize

Serialize

Send Query

Get Result

De-serialize

Tracing the calls: deserialize



So far...

Saw how data flows within **one** connection

How about **concurrent** connections?

How is a connection created?

```
type DB struct { // link  
    // ...  
    openerCh chan struct{}  
    // ...  
}
```

```
func (db *DB) ConnectionOpener() { // link  
    for range db.openerCh {  
        db.openNewConnection()  
    }  
}
```

How is a connection created? (cont.)

```
func (db *DB) maybeOpenNewConnections() { // link  
    // ...  
    for numRequests > 0 {  
        db.numOpen++ // optimistically  
        numRequests--  
        if db.closed {  
            return  
        }  
        db.openerCh <- struct{}{}  
    }  
}
```

How is a connection created? (visualize)

- Modify ``database/sql`` to write logs
- Aggregate events and get counts
- *NOTE: this experiment is not a proper benchmark*

So far...

- database/sql provide app-level connection pool
- pool size is dynamic, but can be controlled by:
 - `SetMaxOpenConns()`
 - `SetMaxIdleConns()`
- Usually: set maximum open conn
 - less chance for runaway connections
 - helps with capacity planning
 - Though it cause other effects (might block requests => increased latency)
- Collect data from `db.Stats()` / source code to plan for capacity

Cleaner code with database/sql/driver

```
func (c CountryType) Value() (driver.Value, error) {  
    return c.String(), nil  
}  
  
func (c *CountryType) Scan(v interface{}) error {  
    val, ok := v.([]byte)  
    if !ok {  
        return fmt.Errorf("invalid data for CountryType: %#v", v)  
    }  
  
    *c = pgCountryTypeToEnum[string(val)]  
    return nil  
}
```

ORM or not ORM?

Basic problems:

- Marshaling data (solved)
- Generating queries (same same, but different)
- Generating queries == generating SQL
- Hard (SQL is expressi) & limit the patterns you can use

What we have not covered

- Performance: prepared query vs one-time
- In-depth connection pooling
- Failure handling

More

- <https://github.com/bradfitz/go-sql-test>
- <https://github.com/DATA-DOG/go-sqlmock>
- <https://docs.google.com/document/d/1F778e7ZSNiSmbju3jsEWzShcb8lIO4kDyfKDNm4PNd8/edit#>
- <https://www.postgresql.org/docs/9.6/static/protocol.html>
- <rtfsc> :D